

COMP20200 Unix Programming

Lecture 7

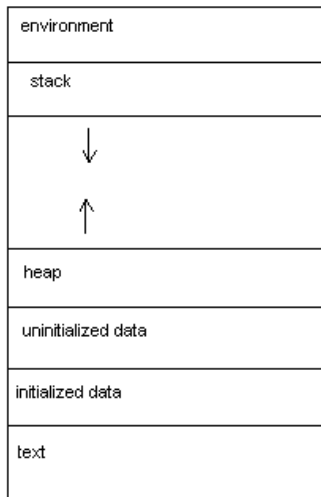
CS, University College Dublin, Ireland



System Calls

- System calls expose functionality of OS to programmes
 - eg. write to hard disk, send & receive data on network.
- Normally program contained within virtual memory space.
 - System call allows program to break out.
- Many system calls, can be platform specific: code isn't portable.
- Standardised POSIX system calls

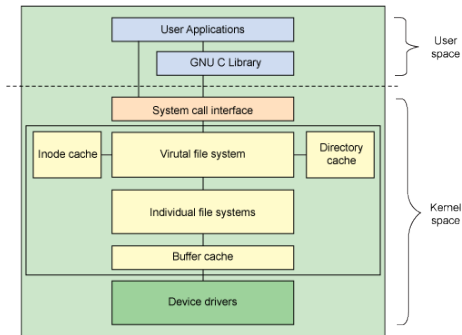

User Space



Virtual memory organization

System Calls

- System calls expose functionality of OS to programmes



System Calls act on:

- Processes
- Files
- Networking Sockets
- Signals
- Inter-process communication
- terminals
- threads
- I/O devices

- System call invoked with a special instruction, written in assembly.
`int 0x80`

- System-call wrappers in standard C library increase portability.

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
int open(const char *pathname, int flags);
```

- C function calls (eg. `fopen()`, `getchar()`, `printf()`) contain several system calls internally.

```
#include <stdio.h>
FILE *fopen(const char *path, const char *mode);
```

- System calls in 2nd Section of manual `man 2 intro`; `man 2 open`.
 - Library functions in 3rd Section of manual `man 3 intro`; `man 3 fopen`

Subset of system calls

- Managing I/O channels: `creat()`, `open()`, `close()`
- Handling I/O operations: `read()`, `write()`
- For random access of files: `lseek()`
- Aliasing and removing files: `link()`, `unlink()`
- Getting file status: `stat()`
- Access control: `access()`, `chmod()`, `chown()`
- Process control: `exec()`, `fork()`, `wait()`, `exit()`
- Process ownership: `getuid()`
- Process ID: `getpid()`
- Process control: `signal()`, `kill()`, `alarm()`
- Changing working directory: `chdir()`
- Manipulate low level memory attributes: `mmap()`, `shmget()`, `mprotect()`, `mlock()`
- Time management: `time()`, `gettimer()`, `settimer()`, `settimeofday()`, `alarm()`
- Creating inter-process communication: `pipe()`

- System call:

```
#include <sys/mman.h>
void *mmap(void *addr, size_t length, int prot,
           int flags, int fd, off_t offset);
int munmap(void *addr, size_t length);
```

- C Library function


```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

- Almost always use function calls

fork()

```
#include <unistd.h>
pid_t fork(void);
```

- Create a new process by duplicating the existing process.
- Existing process becomes the parent,
- newly created process becomes the child process.
- child has a unique PID
- Virtual memory duplicated with copy-on-write (COW)
 - Only when write to memory is physical memory copied.
- Code section also duplicated: same program.
 - Different action based on PID



```
#include <unistd.h> #include <sys/types.h> #include <errno.h>
#include <stdio.h> #include <sys/wait.h> #include <stdlib.h>
```

```
int globalv;
int main(void){
    pid_t childPID;
    int localv = 0;
    childPID = fork();
    if(childPID >= 0) { // fork was successful
        if(childPID == 0) { // child process
            localv++; globalv++;
            printf("Child local:%d global:%d\n", localv, globalv);
        } else { //Parent process
            localv = 10; globalv = 20;
            printf("Parent local:%d global:%d\n", localv, globalv);
        }
    } else { // fork failed
        printf("\n Fork failed.\n");
        return 1; }
    return 0;
}
```

Simple copy with system calls

```
/* cp: copy f1 to f2 */
...
if (argc != 3)
    error("Usage: cp from to");
if ((f1 = open(argv[1], O_RDONLY, 0)) == -1)
    error(...);
if ((f2 = creat(argv[2], PERMS)) == -1)
    error(...);
while ((n = read(f1, buf, BUFSIZ)) > 0)
    if (write(f2, buf, n) != n)
        error(...);
close(f1);
close(f2);
...
```

Note: Full code on next slide

```

#include <stdio.h>
#include <fcntl.h>
#include <unistd.h>
#define PERMS 0666      /* RW for owner, group, others */
void error(char *, ...);

/* cp: copy f1 to f2 */
int main(int argc, char *argv[]) {
    int f1, f2, n;
    char buf[BUFSIZ];
    if (argc != 3)
        error("Usage: cp from to");
    if ((f1 = open(argv[1], O_RDONLY, 0)) == -1)
        error("cp: can't open %s", argv[1]);
    if ((f2 = creat(argv[2], PERMS)) == -1)
        error("cp: can't create %s, mode %03o",
            argv[2], PERMS);
    while ((n = read(f1, buf, BUFSIZ)) > 0)
        if (write(f2, buf, n) != n)
            error("cp: write error on file %s", argv[2]);
    close(f1); close(f2);
    return 0; }

```

- Command line tip(s) of the day:
 - See sorted list of running processes:
 - `top`
 - List all processes:
 - `ps aux`
 - Stop a process
 - `kill <pid>`
 - Force stop of unresponsive: `kill -9 <pid>`
 - Find/kill process by name:
 - `pgrep`
 - `pkill`
- Vi tip(s) of the day:
 - `:sp` Split screen, `:vsp` Vertical split screen
 - `Ctrl+w[jkhl]` or `Ctrl+w Ctrl+w` to move between screens
 - `:tabnew <file name>` Opens new tab.
 - `gt` or `gT` move between tabs
 - `qa` quit all, `wqa` write & quit all