

# COMP20200 Unix Programming

## Lecture 11


CS, University College Dublin, Ireland



# Moors law


- Gordon E. Moore in 1965 predicted “the number of transistors on integrated circuits doubles approximately every two years.”
- Adjusted to doubling performance every 18 months.

Microprocessor Transistor Counts 1971-2011 & Moore's Law



- Need parallelism to use all these transistors: multi-core, many-core.

## Microprocessor Transistor Counts 1971-2011 & Moore's Law



- Blocking IO
  - writing to disk or network
  - Could wait, asynchronous IO or spawn a separate thread
- Server
  - receiving multiple clients requests (eg. web server)
- User Interface
  - Respond to user input while program engine busy
- Performance on machines
  - Desktop, mobile, supercomputer clusters.
  - SMPs (shared memory multiprocessor), GPU-accelerators, clusters
  - Program must be parallel for performance.
  - SIMD (single instruction, multiple data)

# Processes and Threads

- `fork()` returns two completely independent copies of the original process.
  - Each process has own address space, own copies of variables independent of same variables in other process.
  - provides memory protection and therefore stability
  - problem when multiple processes working on same task/problem
- Light-weight process (LWP): POSIX threads or pthreads
  - set of C programming language types, functions and constants
  - Around 100 Pthreads procedures, all prefixed “pthread\_”

# POSIX threads

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>

#define NUM_THREADS      5

void *TaskCode(void *argument) {
    int tid;

    tid = *((int *) argument);
    printf("Hello World! It's me, thread %d!\n", tid);

    /* optionally: insert more useful stuff here */

    return NULL;
}
```

```

int main(void) {
    pthread_t  threads[NUM_THREADS];
    int  thread_args[NUM_THREADS];
    int rc, i;

    /* create all threads */
    for (i=0; i<NUM_THREADS; ++i) {
        thread_args[i] = i;
        printf("In main: creating thread %d\n", i);
        rc = pthread_create(&threads[i], NULL, TaskCode,
                           (void *)&thread_args[i]);
        assert(0 == rc);
    }

    /* wait for all threads to complete */
    for (i=0; i<NUM_THREADS; ++i) {
        rc = pthread_join(threads[i], NULL);
        assert(0 == rc);
    }

    exit(EXIT_SUCCESS);
}

```

# Pthreads

- When compiling with gcc, must add `-pthread` flag.

```
$ gcc -Wall -o pthread_example pthread_example.c -pthread
$
```

Output:

```
In main: creating thread 0
In main: creating thread 1
In main: creating thread 2
Hello World! It's me, thread 1!
Hello World! It's me, thread 0!
In main: creating thread 3
In main: creating thread 4
Hello World! It's me, thread 2!
Hello World! It's me, thread 4!
Hello World! It's me, thread 3!
```



# Race conditions

```
THREAD 1  
a = data;  
a++;  
data = a;
```

```
THREAD 2  
b = data;  
b--;  
data = b;
```

- If above executed serially, no problem, expected result.
- If in parallel it is completely non-deterministic. `data` may be +1, 0, -1.
- Solution: block other threads when writing to `data`
- Pthreads use a data type called *mutex* to achieve this.

# mutex example

THREAD 1

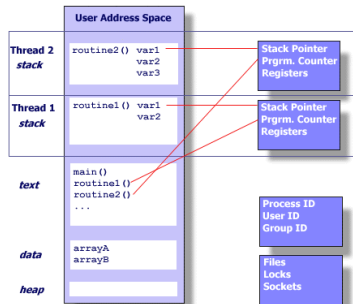
```
pthread_mutex_lock (&mut);  
  
a = data;  
a++;  
data = a;  
pthread_mutex_unlock (&mut);
```

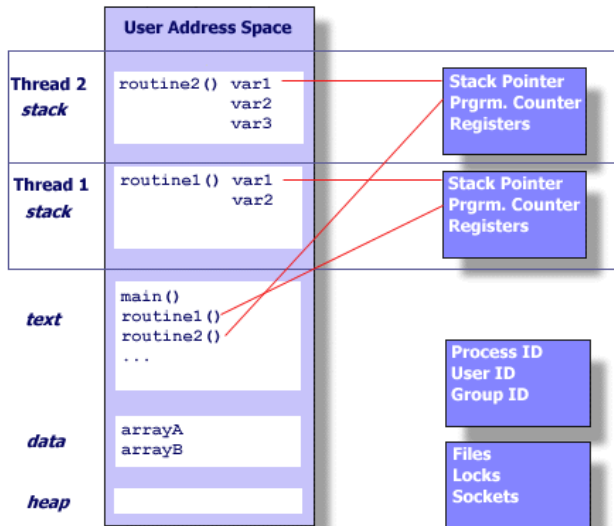
THREAD 2

```
pthread_mutex_lock (&mut);  
/* blocked */  
/* blocked */  
/* blocked */  
/* blocked */  
b = data;  
b--;  
data = b;  
pthread_mutex_unlock (&mut)  
;
```

# POSIX threads

- Independent stream of instructions
- High performance on SMP
- Duplicate only the bare essential resources
- Accomplished by the thread maintaining its own:
  - Stack pointer
  - Registers
  - Scheduling properties (such as policy or priority)
  - Set of pending and blocked signals
  - Thread specific data.





# In UNIX environment, a thread:

- Exists within a process and uses the process resources
- Has its own independent flow of control as long as its parent process exists
- Duplicates only the essential resources it needs to be independently schedulable
- May share the process resources with other threads that act equally independently (and dependently)
- Dies if the parent process dies
- Is "lightweight" because most of the overhead has already been accomplished through the creation of its process.

# Threads share the process resources

Because threads within the same process share resources:

- Changes made by one thread to shared resources will be seen by all threads.
  - eg. closing a file
- Two pointers having the same value point to the same data.
- Reading and writing to the same memory locations is possible
  - but requires explicit synchronization by the programmer.