

# COMP20200 Unix Programming


## Lecture 12

CS, University College Dublin, Ireland




# Designing Multi-threaded Programs

- Challenges affecting all parallel programs:
  - What type of parallel programming model to use?
  - Problem partitioning
  - Load balancing
  - Communications
  - Data dependencies
  - Synchronization and race conditions
  - Program complexity, Programmer effort/costs/time
- Program must be organized into discrete, independent tasks which can execute concurrently



# Shared Memory Model:

- All threads have access to the same global, shared memory
- Threads also have their own private data
- Programmers are responsible for synchronizing access to (protecting) globally shared data.



# Thread-safeness

- If an application can execute multiple threads simultaneously without “clobbering” shared data or creating “race” conditions it is thread safe.
- Eg. Application creates several threads, each makes a call to the same library routine.
  - If this library routine accesses global memory and does not employ synchronization constructs, then it is not thread-safe.

# Creating and Terminating Threads

- `pthread_create (thread,attr,start_routine,arg)`
- `pthread_exit (status)`
- `pthread_cancel (thread)`
- `pthread_attr_init (attr)`
- `pthread_attr_destroy (attr)`

# Creating Threads

- Initially `main()` program comprises a single, default thread.
- All other threads must be explicitly created by the programmer.
- `pthread_create` creates a new thread and makes it executable.
- Can be called any number of times from anywhere within your code.
- Compiling Threaded Programs with `gcc`: `gcc -pthread ...`

# Skeleton example

```
#include <pthread.h>
void* threadfunction(void* arg) {
    int* incoming = (int*)arg;
    (*incoming)++;
    return NULL;
}
int main(void) {
    pthread_t threadID;
    void* exitstatus;
    int value;
    value=42;
    pthread_create(&threadID, NULL, threadfunction, &value);
    pthread_join(threadID, &exitstatus);
    return 0;
}
```

# Returning results from threads

```
void *threadfunction(void *) {  
    int *code;  
    code=malloc(sizeof(int));  
    // Set the value of (*code) — code[0]  
  
    return (void*)code;  
}
```



# Terminating Threads

- Several ways a thread may be terminated:
  - The thread returns normally from its starting routine. It's work is done.
  - The thread makes a call to the `pthread_exit` subroutine - whether its work is done or not.
  - The thread is canceled by another thread via the `pthread_cancel` routine – not recommended
  - The entire process is terminated due to making a call to either the `exec()` or `exit()`
  - If `main()` finishes first, without calling `pthread_exit` explicitly itself
- If `main()` finishes before the threads it spawned, all of the threads will terminate.
  - Explicitly calling `pthread_exit()` will block `main()` until all threads are finished.

# Pthread Creation and Termination


```
#include <pthread.h> #define NUM_THREADS 5

void *PrintHello(void *threadid) {
    long tid;
    tid = *((long*)threadid);
    printf("Hello World! It's me, thread #%ld!\n", tid);
    pthread_exit(NULL); }

int main (int argc, char *argv[]) {
    pthread_t threads[NUM_THREADS];
    int rc; long t;
    for(t=0; t<NUM_THREADS; t++){
        printf("In main: creating thread %ld\n", t);
        //Not good — open to race condition
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)&t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc
                );
            exit(-1);
        } }
    exit(0); }
```

# Joining

- “Joining” is one way to accomplish synchronization between threads
- The `pthread_join()` subroutine blocks the calling thread until the specified threadid thread terminates.
- The programmer is able to obtain the target thread’s termination return status if it was specified in the target thread’s call to `pthread_exit()`.
- A joining thread can match one `pthread_join()` call. It is a logical error to attempt multiple joins on the same thread.
- Two other synchronization methods, mutexes and condition variables, will be discussed later.



# Application of `pthread`: Dot product

- Multiplication of 2 vectors (also known as scalar or inner product)
- $a = [a_1, a_2, \dots, a_n]$  and  $b = [b_1, b_2, \dots, b_n]$   
 $a \cdot b = \sum_{i=1}^n a_i b_i = a_1 b_1 + a_2 b_2 + \dots + a_n b_n$
- In parallel, partition into chunks.
  - Eg. Vector size 1000 and 4 worker threads:
  - Main thread creates the 4 worker threads.
    - Thread 0 computes:  $S_{local} = a_0 b_0 + a_1 b_1 + \dots + a_{250} b_{250}$
    - ...
    - Thread 3 computes:  $S_{local} = a_{750} b_{750} + a_{751} b_{751} + \dots + a_{999} b_{999}$
  - Each worker:
    - calls a mutex lock
    - does  $S_{global} = S_{global} + S_{local}$
    - calls a mutex unlock
  - Main thread joins all worker threads and outputs answer.

# Application of `pthread`s: Dot product

- Code on next 4 slides
- Here is output:

```
Thread 1 did 250 to 499: mysum=250.000000 global sum=250.000000
Thread 0 did 0 to 249: mysum=250.000000 global sum=500.000000
Thread 3 did 750 to 999: mysum=250.000000 global sum=750.000000
Thread 2 did 500 to 749: mysum=250.000000 global sum=1000.000000
Sum = 1000.000000
```

# Dot product 1 of 4

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

typedef struct {
    double    *a;
    double    *b;
    double    sum;
    int       veclen;
} DOTDATA;

#define NUMTHRDS 4
#define VECLEN 250
DOTDATA dotstr;
pthread_t  callThd[NUMTHRDS];
pthread_mutex_t mutexsum;
```

# Dot product 2 of 4

```
void *dotprod(void *arg){
    int i, start, end, len;
    long offset; double mysum, *x, *y;

    offset = *((long*)arg);

    len = dotstr.vecLen;
    start = offset*len;
    end   = start + len;
    x = dotstr.a;
    y = dotstr.b;

    mysum = 0;
    for (i=start; i<end ; i++) {
        mysum += (x[i] * y[i]);
    }

    pthread_mutex_lock (&mutexsum);
    dotstr.sum += mysum;
    printf("Thread %ld did %d to %d:  mysum=%f global sum=%f\n",
        offset, start, end-1, mysum, dotstr.sum);
    pthread_mutex_unlock (&mutexsum);

    pthread_exit(NULL); }
```

# Dot product 3 of 4

```
int main(int argc, char *argv[]){
    long i; double *a, *b;

    a = (double*) malloc (NUMTHRDS*VECLEN*sizeof(double));
    b = (double*) malloc (NUMTHRDS*VECLEN*sizeof(double));

    for (i = 0; i < VECLEN*NUMTHRDS; i++) {
        a[i] = 1;
        b[i] = a[i];
    }

    dotstr.vecLen = VECLEN;
    dotstr.a = a;
    dotstr.b = b;
    dotstr.sum = 0;

    pthread_mutex_init(&mutexsum, NULL);

    pthread_attr_t attr;
    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
```



# Dot product 4 of 4

```
//open to race condition
//for(i=0;i<NUMTHRDS;i++) {
//  pthread_create(&callThd[i], &attr, dotprod, (void *)&i);
//}
long *off;
off=calloc(NUMTHRDS, sizeof(long));
for(i=0;i<NUMTHRDS;i++,off[i]=i) { //safe
    pthread_create(&callThd[i], &attr, dotprod, off+i);
}

pthread_attr_destroy(&attr);

for(i=0;i<NUMTHRDS;i++)
    pthread_join(callThd[i], NULL);

printf("Sum = %f \n", dotstr.sum);
free(a); free(b);
pthread_mutex_destroy(&mutexsum);
pthread_exit(NULL);
}
```