

# COMP20200 Unix Programming

## Lecture 18

Ravi Reddy Manumachu (ravi.manumachu@ucd.ie)

CS, University College Dublin, Ireland

12/04/2023



- Overview of UNIX interprocess communication (IPC).
- Learn the oldest method of IPC, UNIX pipes.

# UNIX Interprocess communication (IPC)

IPC are tools that processes and threads can use to communicate with one another and to synchronize their actions.

Communication	Synchronization	Signal
<b>pipe, FIFO, stream socket</b>  <b>pseudoterminal</b>  <b>System V message queue, POSIX message queue, datagram socket</b>  <b>System V shared memory POSIX shared memory Memory mapping - anonymous Memory mapping - mapped file</b>	<b>System V semaphore POSIX semaphore - named POSIX semaphore - unnamed</b>  <b>eventfd</b>  <b>file lock (flock()) record lock (fcntl())</b>  <b>futex mutex (threads) condition variable (threads) barrier</b>	<b>standard signal realtime signal</b>
<b>Covered in this Module</b>		
<b>stream and datagram sockets (Lectures 14-17), mutex (L10), pipe (L18)</b>		

Figure: UNIX IPC tools.

# UNIX Interprocess communication (IPC)




Figure: UNIX IPC features.

There are three broad functional categories in IPC:

- *Communication*: Tools concerned with exchanging data between processes.
- *Synchronization*: Tools concerned with synchronizing the actions of processes or threads.
- *Signals*: Can be used for synchronization or communication.  
*Review lecture on signals.*

# IPC communication tools - Data transfer





Figure: IPC data transfer facility.

There are two main categories under *communication*.

The first category is **Data transfer**.

- One process writes to the IPC facility, and another process reads the data.
- Two data transfers required between user memory and kernel memory.
- One transfer from user memory to kernel memory during writing.
- One transfer from kernel memory to user memory during reading.
- Synchronization between the reader and writer processes is automatic.

# IPC communication tools - Shared memory




**Figure:** IPC shared memory facility.

The second category is **Shared memory**.

- Processes exchange information by placing it in a region of physical memory that is shared between the processes.
- The kernel does this by making page-table entries in each process point to the same pages of RAM.
- Communication doesn't require system calls or data transfer between user memory and kernel memory.
- So, this IPC can be fast but synchronizing reads and writes can be tricky.

# IPC data transfer tools - Byte stream




**Figure:** IPC data transfer facility, byte stream.

Following are the main categories under data transfer:

## Byte stream:

- Data exchanged via stream sockets, pipes, and FIFOs is an undelimited byte stream.
- There is no concept of message or message boundaries.
- *We cover stream sockets and pipes in this module.*

# IPC data transfer tools - Message




**Figure:** IPC data transfer facility, message.

## Message:

- Data exchanged via datagram sockets and message queues are delimited messages.
- Each read operation reads a whole message, as written by the writer process.
- It is not possible to read part of a message.
- It is not possible to read multiple messages in a single read operation.



# IPC data transfer tools - Pseudoterminal



**Figure:** IPC data transfer facility, pseudoterminal.

## Pseudoterminal:

- It is a pair of virtual devices, known as the master and slave.
- This pair of devices provides an IPC channel for data transfer between the two devices.
- Allows a user on one host to operate a terminal-oriented program on another host connected via a network.
- Employed in terminal emulators and *ssh*, which provides network login service.

Following are the main categories under synchronization:

## Semaphores:

- A semaphore is a kernel-maintained integer whose value is never permitted to fall below 0.
- A process can increase or decrease the value of the semaphore.
- If an attempt is made to decrease the value of the semaphore below 0, then the kernel blocks the operation.
- A process decrements a semaphore in order to reserve exclusive access to some shared resource, and
- After completing work on the resource, a process increments the semaphore so that the shared resource is released for use.

## File locks:

- Synchronize the actions of multiple processes operating on the same file.
- *fcntl()* system call provides record locking, allowing processes to place multiple read and write locks on different regions of the same file.

## Mutexes and Condition variables:

- A mutex (short for mutual exclusion) ensures that only one thread at a time can access a shared variable.
- A mutex prevents multiple threads from accessing a shared variable at the same time.
- A condition variable allows one thread to inform other threads about changes in the condition/state of a shared variable, and
- Allows the other threads to wait (block) for such notification.
- Revisit lecture on threads (L11).

# UNIX IPC using Pipes

# IPC: Data transfer using pipes

Pipes can be used to pass data between *related* processes.

Consider the following:

```
$ cat iserver.c | wc -l
```

- *cat* prints the contents of the file *iserver.c* to standard output.
- *wc -l* outputs the numbers of lines in the input given to it.
- Shell creates two processes (using *fork()* and *exec()*).
- One process executes *cat* and the other, *wc*.
- The output produced by *cat* process is piped as input to *wc*.
- The vertical bar between the two programs represents a pipe.

## IPC Pipe

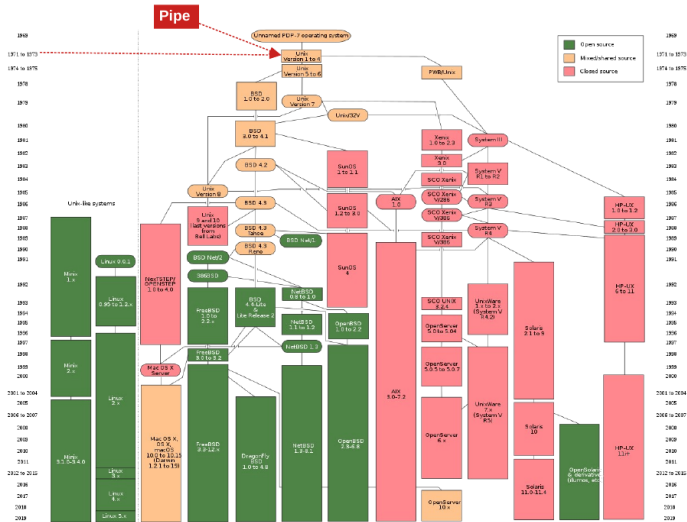


Figure: Pipes appeared in Second Edition UNIX in 1972.

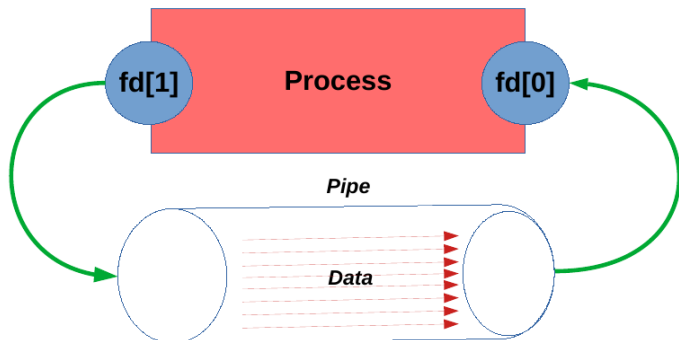
# IPC Pipe API

```
#include <unistd.h>
int pipe(int fd[2]);
```

- Two open file descriptors are returned in *fd*, read end in *fd[0]* and the write end in *fd[1]*.
- The *read()* and *write()* system calls can be used to perform I/O on the pipe.
- Pipes are used for communication between related processes. For example:
  - Parent process and child process.
  - Parent process and its grandchild.
  - Two sibling processes (two child processes created by a parent).




# IPC Pipe



**Figure:** File descriptors after creation and direction of data flow.

- A pipe is a byte stream. There is no concept of messages or message boundaries when using a pipe.
- Pipes are unidirectional. One end of the pipe is used for writing, and the other end is used for reading.
- Pipes have a limited capacity. A pipe is a buffer maintained in kernel memory.
- However, the application doesn't need to know the capacity (or maximum) of a pipe.
- Once a pipe is full, further writes to the pipe will block until the reader removes some data from the pipe.

# IPC Pipe - Parent to child



**Figure:** Simple pipe program creating a channel from parent to child.

# Simple pipe program

```
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>
#define BUFSIZE 10
int main(int argc, char *argv[])
{
    int pfd[2];
    char buf[BUFSIZE];

    if (pipe(pfd) == -1)
        exit(EXIT_FAILURE);
    /* Cont'd... */
}
```

- The program takes a string as an input. The parent process writes the string to be received by the child process.
- *pipe()* system call is invoked to create the pipe.

# Simple pipe program: Child process code

```
switch (fork()) {  
case 0:  
    if (close(pfd[1]) == -1)  
        exit(EXIT_FAILURE);  
/* Cont'd... */
```

The child process closes the write end of the pipe (*pfd[1]*) since it is required to read only from the parent process.

# Simple pipe program: Child process code

```
switch (fork()) {  
case 0:  
    for (;;) {  
        numRead = read(pfd[0], buf, BUFSIZE);  
        if (numRead == -1)  
            exit(EXIT_FAILURE);  
        if (numRead == 0)  
            break;  
        /* Cont'd... */  
    }  
}
```

- The child enters a loop to read blocks of data (of size BUFSIZE) and writes it to the standard output.
- *numRead* of 0 represents end-of-file (EOF) condition.
- The child process exits the loop when it receives the end-of-file (EOF) condition.

# Simple pipe program: Child process code

```
switch (fork()) {  
case 0:  
  
    if (write(1, buf, numRead) != numRead)  
        exit(EXIT_FAILURE);  
} /* for loop end */  
  
write(1, "\n", 1);  
/* Cont'd... */
```

- The child process writes what it has read to the standard output.
- Note the use of *write()* system call to write to standard output represented by the file descriptor 1.

# Simple pipe program: Child process code

```
switch (fork()) {  
case 0:  
    if (close(pfd[0]) == -1)  
        exit(EXIT_FAILURE);  
  
    _exit(EXIT_SUCCESS);  
  
/* Cont'd... */
```

- It then closes the read-end of the pipe (*pfd[0]*).
- A child process terminates by issuing an *\_exit()* call.



# Simple pipe program: Parent process code

```
switch (fork()) {  
default:  
    if (close(pfd[0]) == -1)  
        exit(EXIT_FAILURE);  
    if (write(pfd[1], argv[1], strlen(argv[1])) != strlen(argv[1]))  
        exit(EXIT_FAILURE);  
    ...  
/* Cont'd */
```

- The parent process closes the read end of the pipe (*pfd[0]*).
- The string to be written to the pipe is passed to the program as the first argument (*argv[1]*).
- It then writes the input string to the write end of the pipe, *pfd[1]*.

# Simple pipe program: Parent process code

```
switch (fork()) {  
default:  
    if (close(pfd[1]) == -1)  
        exit(EXIT_FAILURE);  
    wait(NULL);  
    exit(EXIT_SUCCESS);  
}
```

- The parent process then closes the write end of the pipe so that the child process encounters EOF condition.
- It finally waits for the child to finish before exiting.
- The *wait()* call blocks until the child terminates.
- *What happens if child process terminates before parent process calls wait()?*

# Simple pipe program: Parent process code

```
switch (fork()) {  
default:  
    if (close(pfd[1]) == -1)  
        exit(EXIT_FAILURE);  
    wait(NULL);  
    exit(EXIT_SUCCESS);  
}
```

- As discussed in the previous lecture, *wait()* call allows parent to reap the child process and disallow it to become a zombie.
- The parent passes NULL to the *wait()* call since it is not interested in the exit status of the child.
- *Why no SIGCHLD signal handler?*

# exit() and \_exit() calls

- *exit()* is a library function; *\_exit()* is a system call.
- Programs typically call the *exit()* library function.
- *exit()* performs the following actions:
  - Exit handlers.
  - stdio stream buffers are flushed.
  - *\_exit()* system call is invoked.
- Programs also exit by issuing **return n**, which is the same as *exit(n)*.

## `exit()` and `_exit()` calls

- In an application with parent and child processes, typically only one of the parent and child terminate by calling `exit()`.
- The other should terminate by calling `_exit()`.
- This is done so that only one process calls exit handlers and flushes *stdio* buffers.
- When this is not done, you observe duplication of `printf` strings to the standard output.

# Simple pipe program: Execution

spipe.c - Contains the simple pipe program code.

```
$ gcc -o spipe spipe.c
```

```
$ ./spipe "UNIX was initially developed at Bell Labs and became  
operational on a PDP-7 in 1970."
```

```
Parent: Writing 'UNIX was initially developed at Bell Labs and  
became operational on a PDP-7 in 1970.' to child.
```

```
Child received
```

```
'UNIX was initially developed at Bell Labs and became operational  
on a PDP-7 in 1970.'
```

# Lectures 19-21: Bash scripting language

We will focus on Bash scripting language in the next three lectures.

Q & A